

---

# Generative Adversarial Networks (GANs)

— Daniel Foronda-Pascual —

---

# Applications

- Image generation
- Style transfer
- Data augmentation
- Super-resolution
- Image-to-image translation
- Inpainting & restoration
- Text-to-image synthesis
- Synthetic medical data
- Deepfakes & facial reenactment

## Image generation



2014



2015



2016



2017



2018

# Image to image translation

Labels to Street Scene




input



output

Aerial to Map




input



output

Labels to Facade



input



output

Day to Night



input



output

BW to Color



input



output

Edges to Photo



input




output

# Style transfer





## Super resolution



SRGAN

ESRGAN

Ground Truth

## text2image

The small bird has a red head with feathers that fade from red to gray from head to tail

Stage-I  
images



Stage-II  
images



This bird is black with green and has a very short beak

Stage-I  
images




Stage-II  
images



# Facial rejuvenation and aging



(a)



(b)




(c)



(d)




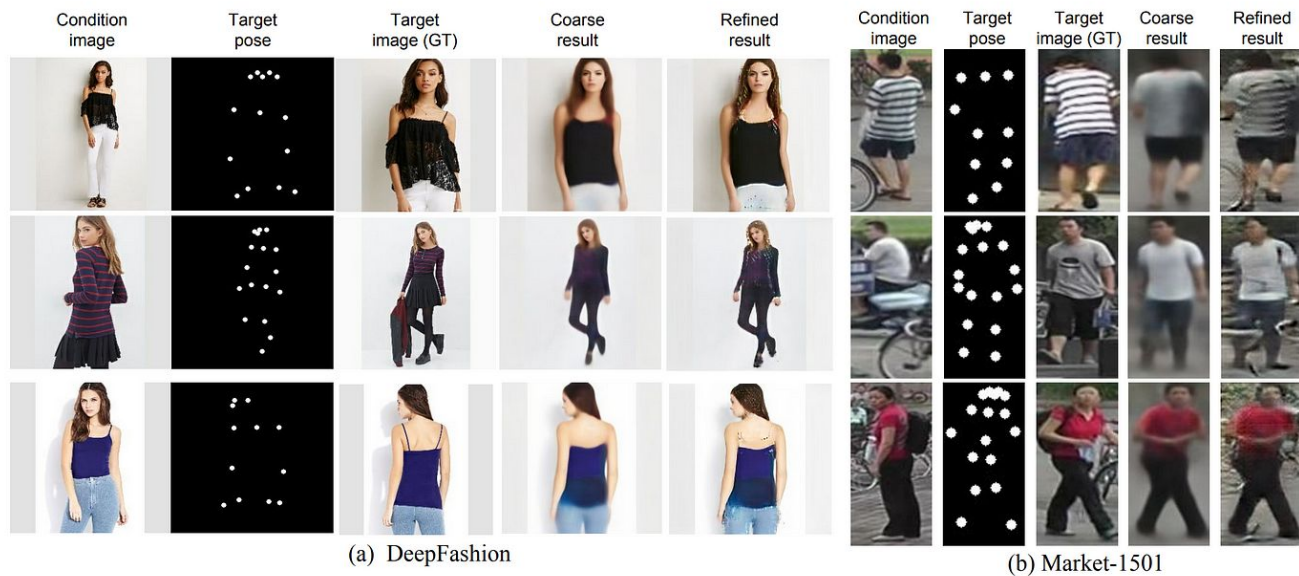
## Fill missing parts of the image (inpainting)



(a) Face Editing

(b) Object Removal

# New human poses generation




# Great-grandfather of GANs

**HERACLITUS OF EPHESUS (c. 535 – 480 BCE)**

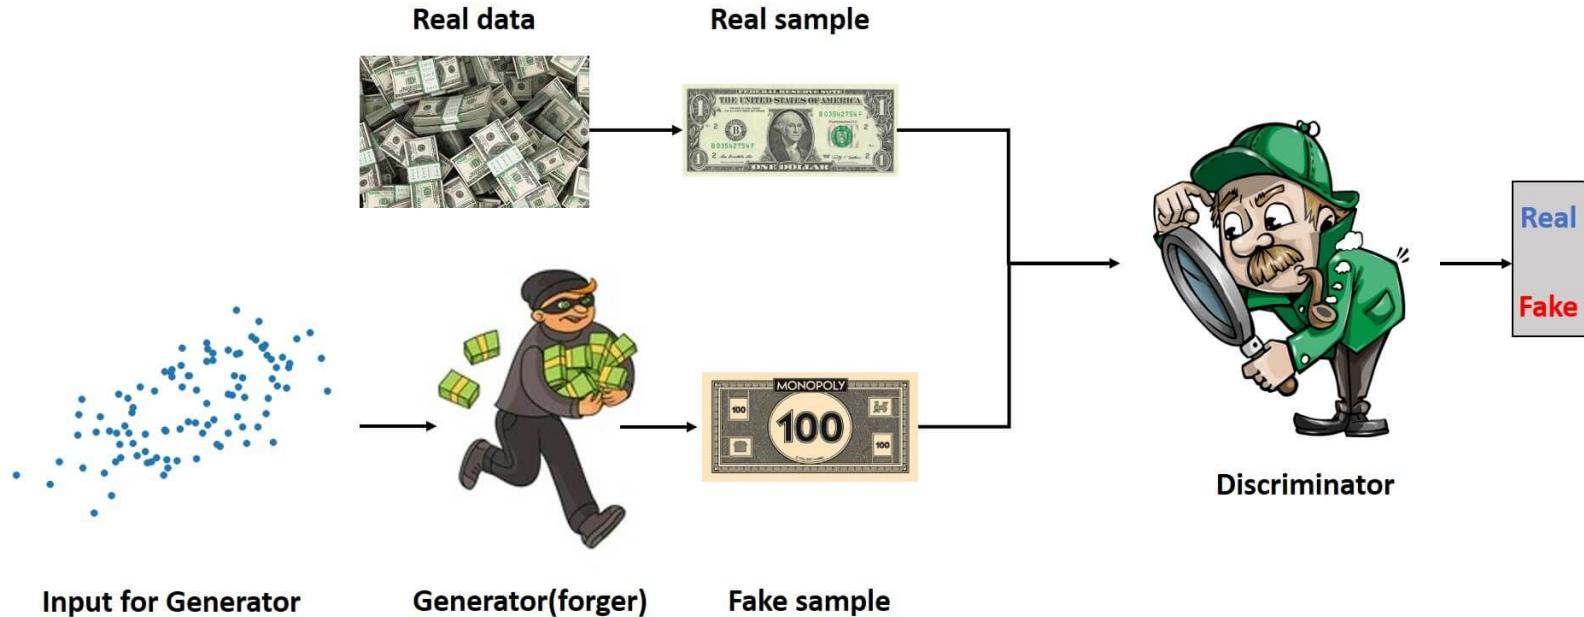
Opposition is the father of all things, the king of all.

Opposition brings concord. Out of discord comes the  
fairest harmony.

Everything changes, nothing remains the same.

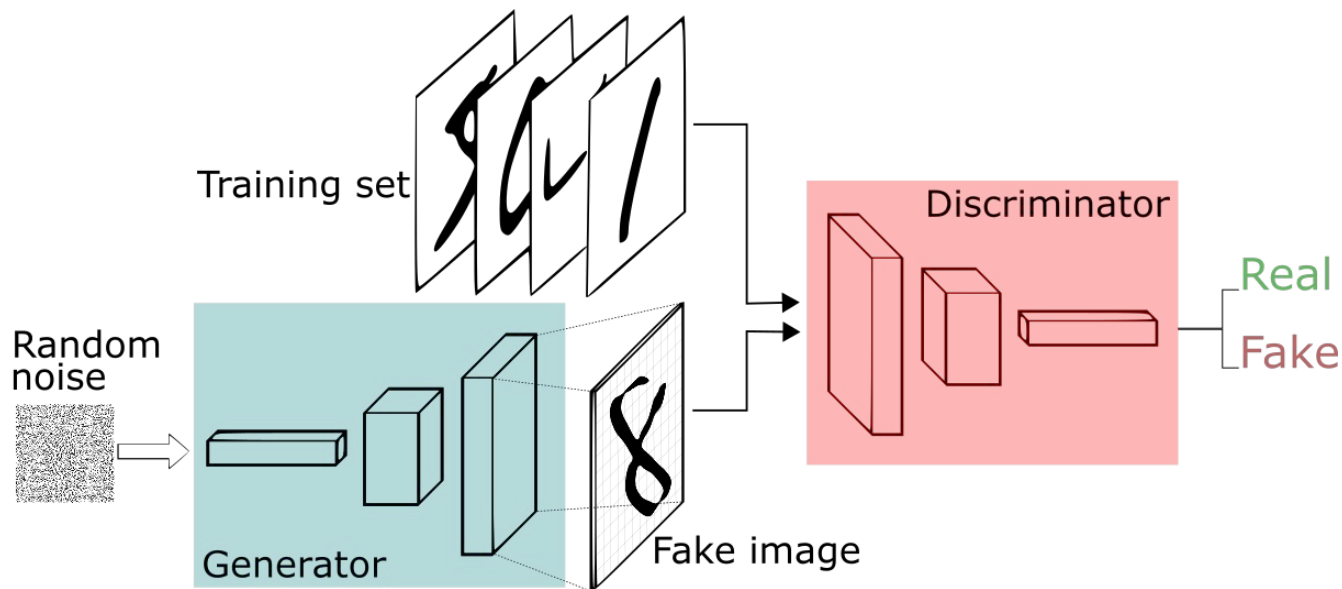


# How GANs work - Idea





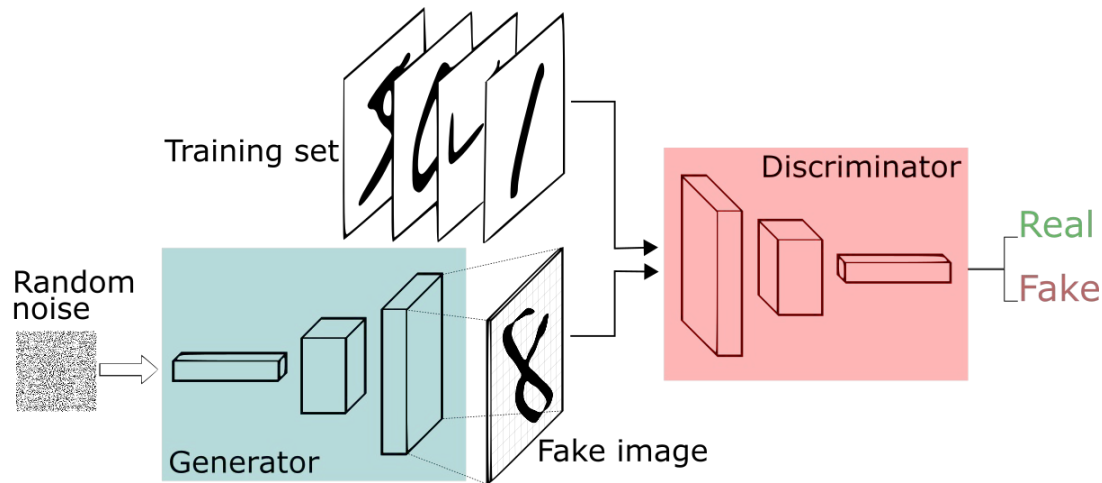
# How GANs work - Idea



# How GANs are trained - The minimax game

**Generator:**  $\frac{1}{m} \sum_{i=1}^m \log(1 - D(G(z^{(i)})))$  *Minimize*

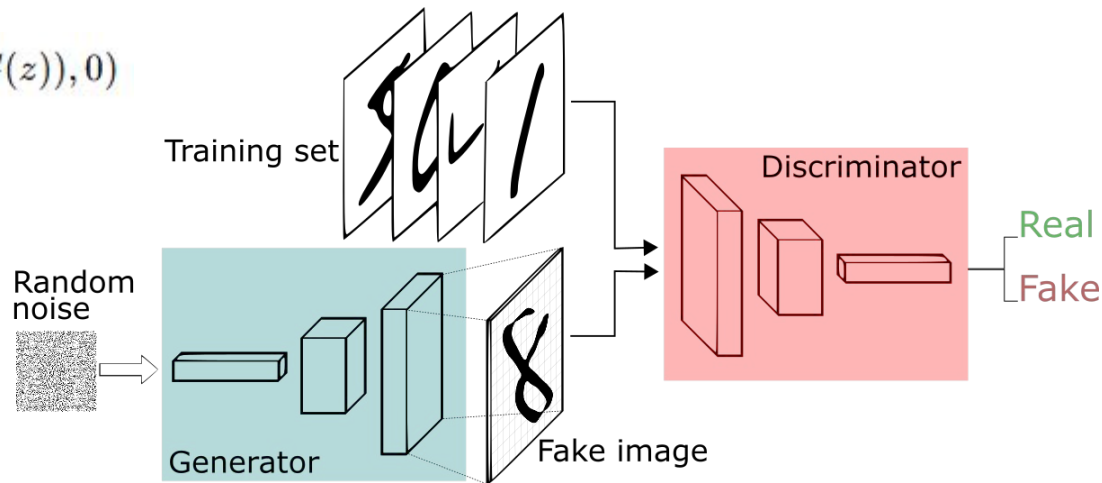
$$\mathcal{L}_G = \text{BCE}(D(G(z)), 1)$$



# How GANs are trained - The minimax game

**Discriminator:**  $\frac{1}{m} \sum_{i=1}^m [\log D(x^{(i)}) + \log(1 - D(G(z^{(i)})))]$  *Maximize*

$$\mathcal{L}_D = \text{BCE}(D(x), 1) + \text{BCE}(D(G(z)), 0)$$



# How GANs are trained - The minimax game

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

The Nash equilibrium of this particular game is achieved at:

- $P_{\text{data}}(x) = P_{\text{gen}}(x) \quad \forall x$
- $D(x) = \frac{1}{2} \quad \forall x$



# Common training challenges

- **Mode collapse:** *Generator outputs lack diversity.*  
→ Use minibatch discrimination, feature matching, WGAN
- **Vanishing gradients:** *Discriminator becomes too strong, generator stops learning.*  
→ WGAN, label smoothing, or modify the loss.
- **Unstable training:** *Training fails to converge or oscillates.*  
→ Use spectral normalization, gradient penalties, and tune hyperparameters.

# Brief history (2014 - 2021)

## ♦ 2014

Ian Goodfellow et al. introduce GANs in “*Generative Adversarial Nets*” → breakthrough in unsupervised learning.

Goodfellow, I. et al. (2014). *Generative Adversarial Nets*. NeurIPS. <https://arxiv.org/abs/1406.2661>

## ♦ 2015–2016

**DCGAN** improves image quality.

Radford, A., Metz, L., & Chintala, S. (2016). *Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks*. ICLR.

<https://arxiv.org/abs/1511.06434>

**cGAN** allows controlled generation.

Mirza, M., & Osindero, S. (2014). *Conditional Generative Adversarial Nets*. arXiv preprint. <https://arxiv.org/abs/1411.1784>

## ♦ 2017

**CycleGAN** enables image-to-image translation without paired data.

Zhu, J.-Y., Park, T., Isola, P., & Efros, A.A. (2017). *Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks*. ICCV.

<https://arxiv.org/abs/1703.10593>

**Wasserstein GAN (WGAN)** introduces a more stable training method.

Arjovsky, M., Chintala, S., & Bottou, L. (2017). *Wasserstein GAN*. ICML. <https://arxiv.org/abs/1701.07875>

## ♦ 2018–2021

**StyleGAN / StyleGAN2 / StyleGAN3** raise the bar in photo-realistic synthesis.

StyleGAN (2019): Karras, T. et al. *A Style-Based Generator Architecture for GANs*. CVPR. <https://arxiv.org/abs/1812.04948>

StyleGAN2 (2020): Karras, T. et al. *Analyzing and Improving the Image Quality of StyleGAN*. CVPR. <https://arxiv.org/abs/1912.04958>

StyleGAN3 (2021): Karras, T. et al. *Alias-Free Generative Adversarial Networks*. NeurIPS. <https://arxiv.org/abs/2106.12423>

**BigGAN** enables high-res generation at scale.

Brock, A., Donahue, J., & Simonyan, K. (2019). *Large Scale GAN Training for High Fidelity Natural Image Synthesis*. ICLR. <https://arxiv.org/abs/1809.11096>

# DCGAN (Deep Convolutional GAN)

Radford, A., Metz, L., & Chintala, S. (2016). Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks. ICLR. <https://arxiv.org/abs/1511.06434>


**Key idea:** Use CNN architectures for both Generator and Discriminator

## Main Contributions:

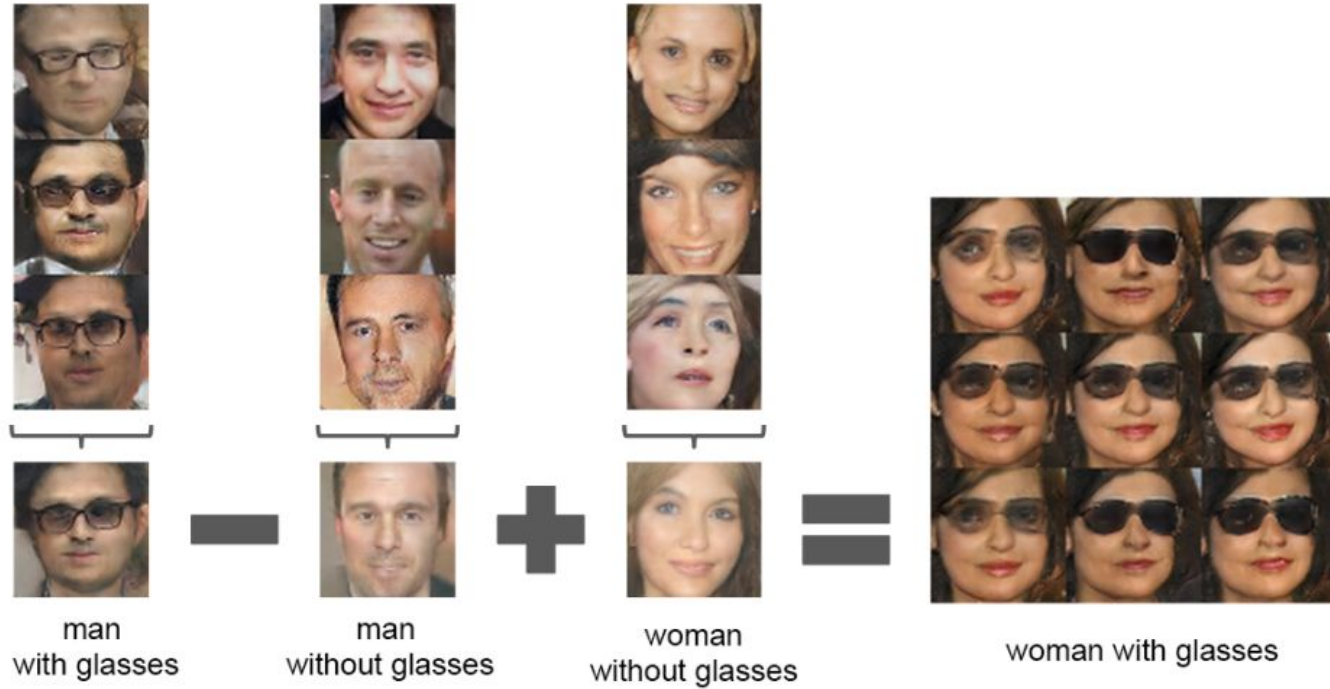
- Replace fully connected layers with convolutions / transposed convolutions
- Use batch normalization for stable training
- Apply ReLU in Generator, LeakyReLU in Discriminator
- Remove pooling layers → let strides do the down/up-sampling

## Advantages:

- More stable training
- Better image quality
- Simple and effective architecture for image generation tasks



# DCGAN (Deep Convolutional GAN)





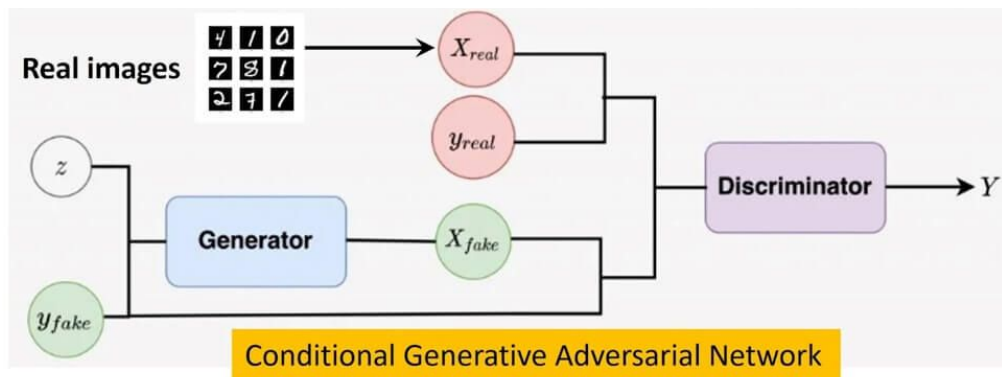
# cGAN (conditional GAN)

📄 Mirza, M., & Osindero, S. (2014). Conditional Generative Adversarial Nets. arXiv preprint. <https://arxiv.org/abs/1411.1784>


**Key idea:** Add conditioning information (e.g., class labels) to both Generator and Discriminator

## How it works

- Generator receives input noise  $z$  and label  $y \rightarrow$  generates image  $G(z | y)$
- Discriminator evaluates whether an image is real/fake given label  $y \rightarrow D(x | y)$



# CycleGAN

 Zhu, J.-Y., Park, T., Isola, P., & Efros, A.A. (2017). Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks. ICCV.  
<https://arxiv.org/abs/1703.10593>

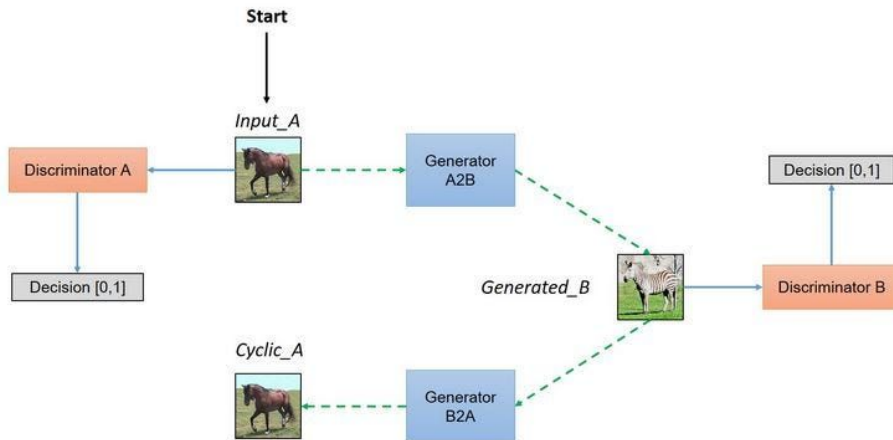
**Key idea:** Learn image-to-image translation between two domains **without paired data**

## Losses:

- Adversarial loss: Two GAN losses (one for each translation direction) encourage the generated images to be indistinguishable from real images in the target domain.
- Cycle-consistency loss: Ensures that translating an image to the other domain and back reconstructs the original (e.g.,  $A \rightarrow B \rightarrow A \approx A$ ). This enforces content preservation.
- Identity loss (optional): Penalizes changes when the input image already belongs to the target domain, helping preserve colors and structure.

## Applications:

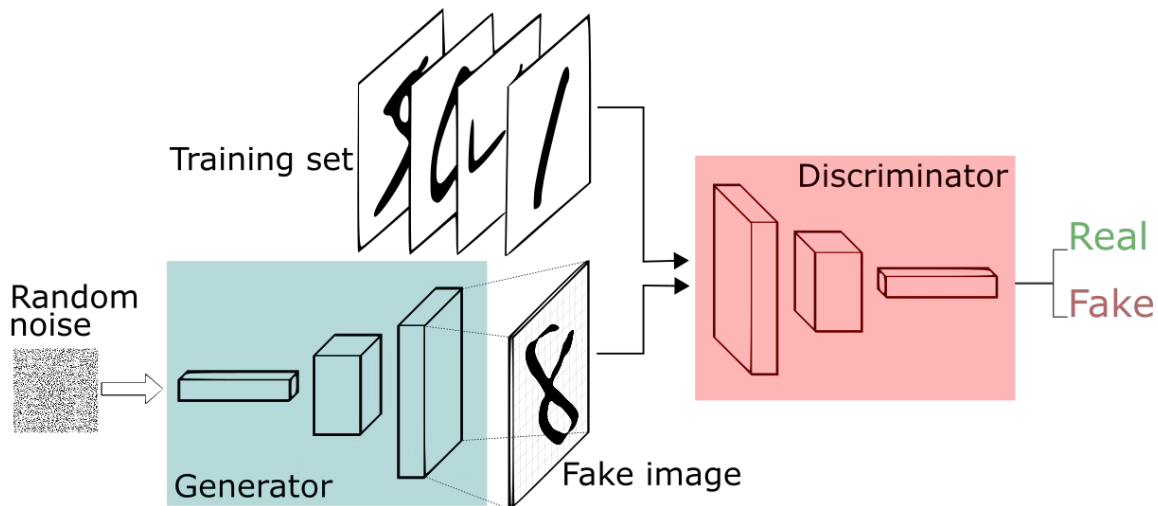
- Artistic style transfer (painting  $\leftrightarrow$  photo)
- Object translation (apples  $\leftrightarrow$  oranges)
- Medical image translation (MRI  $\leftrightarrow$  CT)



# WGAN (Wasserstein GAN)


Arjovsky, M., Chintala, S., & Bottou, L. (2017). Wasserstein GAN. ICML. <https://arxiv.org/abs/1701.07875>

**Problem:** When the discriminator gets too good, it stops giving useful feedback to the generator—gradients vanish, learning slows down, and training becomes unstable.



# WGAN (Wasserstein GAN)

**Problem:** Vanishing gradients when Generator is not good



$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$



# WGAN (Wasserstein GAN)

**Wasserstein distance** (Earth mover's distance): It measures the minimum effort required to transform one distribution into another




$$W(p_{\text{data}}, p_G) = \inf_{\gamma} E_{(x,y) \sim \gamma(x,y)} [\|x - y\|]$$

$\gamma(x, y)$  is a *distribution* over  $x, y$   
with marginals  $\gamma_X(x) = p_{\text{data}}(x)$  and  $\gamma_Y(y) = p_G(y)$

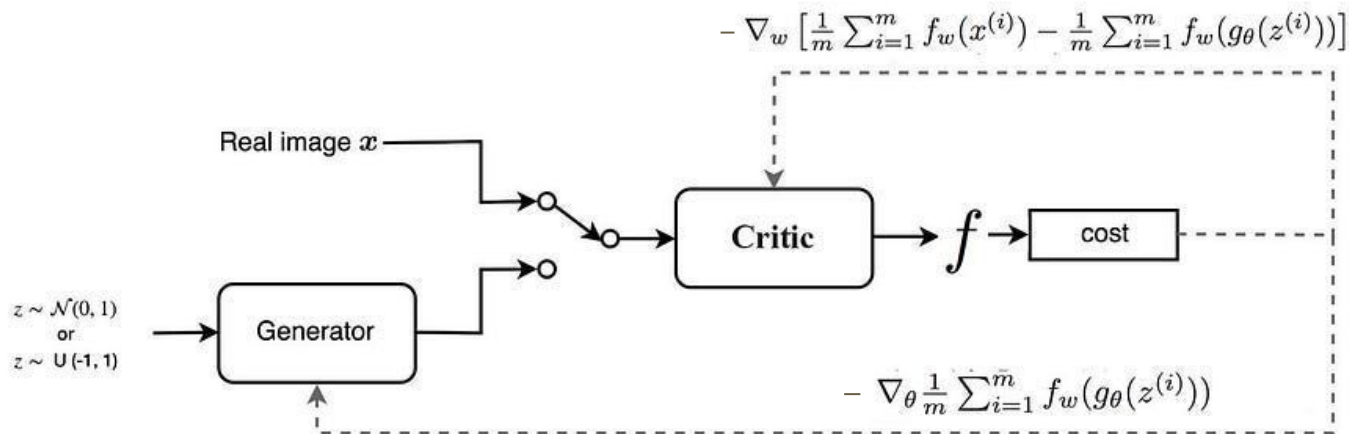
dualidad de Kantorovich-Rubinstein

$$W(p_{\text{data}}, p_G) = \sup_{\|f\|_L \leq 1} E_{p_{\text{data}}}[f(x)] - E_{p_G}[f(x)]$$



# WGAN (Wasserstein GAN)

**Discriminator now is called Critic:** The critic assigns higher scores to more realistic images and lower scores to fake ones, reflecting their distance from the real data distribution.



# 1-Lipschitz Constraint in WGANs

To approximate the Wasserstein distance correctly, the critic must be a 1-Lipschitz function. Different techniques have been proposed to enforce this:

- **Weight Clipping** (WGAN, 2017)

Clip critic weights to a fixed range (e.g.,  $[-0.01, 0.01]$ )

Simple but can lead to capacity issues and poor gradients

- **Gradient Penalty** (WGAN-GP, 2017)

Add a penalty term to the loss to enforce that the gradient norm is close to 1:

$$\lambda \cdot (\|\nabla_{\hat{x}} f(\hat{x})\|_2 - 1)^2$$

where  $\hat{x}$  is interpolated between real and fake samples

More stable and widely used

- **Spectral Normalization** (2018)

Normalize each layer's weight matrix by its largest singular value (spectral norm)

Lightweight and effective; used in many modern GANs

```
# Get scores for real and fake images
real_scores = critic(real_images)
fake_scores = critic(fake_images)

# Get gradients
loss_critic = fake_scores.mean() - real_scores.mean()
gradient_penalty = ((gradient_norm - 1) ** 2).mean() * lambda_gp

# Add the gradient penalty to the critic loss and backpropagate
loss_critic_total = loss_critic + gradient_penalty
loss_critic_total.backward()
```

# GANs in 2025

Today GANs are no longer dominant in image generation.

## Where GANs remain useful:

- ⚡ Fast inference for real-time applications
- 🧠 Low-data regimes where training data is limited
- 🔧 Custom setups with tight resource constraints
- 🎨 Creative tools in art, design, and music

## ⚠️ Challenges remain:

Training instability

Mode collapse

Difficulty scaling to complex data

## Competing with:

Diffusion Models → better quality & stability

Transformers → better semantic control

Latent models → efficient + high-fidelity

- More than **55,000 people** have been reported **killed**.
- **80% of Palestinians killed are civilians**.
- **70% of the Palestinians killed in residential buildings or similar housing were women and children**.

Wikipedia







# Jupyter notebook

<https://www.kaggle.com/code/rafat97/pytorch-wasserstein-gan-wgan>